

1

File system calls

- Unix system calls
 - => call directly to the kernel functions of the OS
 - (the lower level of file processing)
- Standard functions in C library:
 - open, fprintf, fscanf, fclose, fseek, etc**
 - (the higher level of file processing)

2

The *open* system call

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int open(const char *pathname, int flags [,mode_t mode]);

flags: O_RDONLY, O_WRONLY, O_RDWR, O_NONBLOCK, O_APPEND, O_CREAT, O_EXCL (with O_CREAT), O_TRUNC
mode:
  * access rights, works in accordance with umask
  * S_IRUSR, S_IWUSR, S_IRGRP, S_IWGRP, S_IROTH, S_IWOTH, S_IXOTH

Return:
  -1 - error
  >0 - the file descriptor
```

man open

3

The *open* system call. Example 1

```
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>

//example open call
int main() {
    int f1 = open("test1", O_CREAT);
    int f2 = open("test2", O_TRUNC);
    int f3 = open("test3", O_CREAT, S_IRUSR | S_IWUSR | S_IXUSR);
    int f4 = open("test4", O_CREAT, S_IRUSR, S_IWUSR, S_IXUSR);
    int f5 = open("test5", O_WRONLY);
    return 0;
}
```

Example: Test Program 2018

4

The *open* system call. Example 2

```
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <stdio.h>

//example open call
int main() {
    int f1 = open("cont1", O_WRONLY);
    int f2 = open("cont1", O_RDONLY);
    printf("fd and fd", f1, f2);
    return 0;
}
```

Example: Test Program 2018

5

The *open* system call. Example 3

```
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <stdio.h>

//example open call with creat command

int main() {
    printf("file", open("f1", O_CREAT, S_IRWXU) == -1 ? "Error: No error" : "");
    printf("file", open("f2", O_CREAT | O_EXCL, S_IRWXU) == -1 ? "Error: No error" : "");
    return 0;
}
```

Example: Test Program 2018

6

The *creat* system call

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int creat(const char *pathname, mode_t mode);

/* equivalent to specifying the options: O_WRONLY|O_CREAT|O_TRUNC for the open function */

man creat
```

7

The *creat* system call. Example

```
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <stdio.h>

//example creat call

int main() {
    printf("file", creat("f1", S_IRWXU) == -1 ? "Error: No error" : "");
    return 0;
}
```

Example: Test Program 2018

8

The *close* system call

```
#include <unistd.h>

int close (int fd);

/* fd - file descriptor obtained by open */

man close
```

9

Functions *read* and *write*

```
#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count);
ssize_t write(int fd, void *buf, size_t count);

Return:
  0 - EOF
  -1 - error
  >0 - number of bytes operated.

man 2 read
man 2 write
```

10

Functions *read* and *write*. Example 1

```
#include <fcntl.h>
...
#include <stdio.h>

// read write example
int main() {
    int f1 = open("cont1", O_RDONLY);
    char buff[256];
    read(f1, buff, 7);
    write(1, buff, 7);
    return 0;
}
```

Example: Test Program 2018

11

Functions *read* and *write*. Example 2

```
...
//example of reading a sentence and adding a word
int main() {
    int f1 = open("cont1", O_RDONLY);
    char buff[256];
    read(f1, buff, 1);
    int i = 0;
    while(buff[i++] != '\n') read(f1, buff + i, 1);
    buff[i] = 0;
    write(1, buff, 1);
    write(f1, "end", 3);
    return 0;
}
```

Example: Test Program 2018

12

Functions *read* and *write*. Example 3

```
...
#include <stdio.h>

// read write example
int main() {
    int f1 = open("cont1", O_WRONLY | O_RDONLY); // not equivalent to O_RDWR
    char buff[256];
    buff[0] = 'A';
    printf("file", read(f1, buff, 1) == -1 ? "Error: No error" : "");
    printf("file", write(f1, buff, 1) == -1 ? "Error: No error" : "");
    return 0;
}
```

Example: Test Program 2018

13

Function *lseek*

```
#include <sys/types.h>
#include <unistd.h>

off_t lseek(int fd, off_t offset, int whence);

offset: the number of bytes on which the move is made;
whence: position relative to:
  SEEK_SET - the beginning of the file;
  SEEK_CUR - current position;
  SEEK_END - end of file.

man lseek
```

14

Function *lseek*. Example 1

```
// lseek example
int main() {
    int f1 = open("cont1", O_RDONLY);

    char buff[1];
    lseek(f1, 7, SEEK_CUR);
    read(f1, buff, 1);
    printf("io", *buff);

    lseek(f1, -1, SEEK_CUR);
    read(f1, buff, 1);
    printf("io", *buff);

    lseek(f1, 0, SEEK_SET);
    read(f1, buff, 1);
    printf("io", *buff);
}
```

Example: Test Program 2018

15

Function *lseek*. Example 2

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <stdio.h>

int main() {
    int f1 = open("cont", O_APPEND);
    printf("file", lseek(f1, 0, SEEK_SET));
    printf("file", write(f1, "A", 1));
    return 0;
}
```

Example: Test Program 2018

16

Other examples with system calls

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

# ...
int fd1, fd2;

fd1 = open ("alina.txt", O_WRONLY | O_TRUNC);
if (fd1 < 0) {
    perror("open");
    exit (EXIT_FAILURE);
}

fd2 = open ("dan.txt", O_RDWR | O_CREAT, 0644);
# ...
```

17

Complete example

```
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

int main (void) {
    int fd;
    char *buff;
    ssize_t n;

    /* we open the file */
    fd = open ("abi.txt", O_RDONLY);
    if (fd < 0) {
        perror ("open");
    }
}
```

18

Standard functions in the C library

```
#include <stdio.h>

FILE *fopen(const char *path, const char *mode);
int fclose(FILE *fp);

int fprintf(FILE *stream, const char *format, ...);
int fscanf(FILE *stream, const char *format, ...);
size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream);
size_t fwrite(void *ptr, size_t size, size_t nmemb, FILE *stream);
int fseek(FILE *stream, long offset, int whence);

man 3.
```

19

System calls *dup* and *dup2*

```
#include <unistd.h>

int dup(int oldfd);
int dup2(int oldfd, int newfd);

Duplicates the file descriptor oldfd as follows:
  - oldfd and newfd refer to the same physical file;
  - the access mode is retained from the opening;
  - both descriptors share the same current pointer in the file.

Return:
  -1 - error;
  >0 - the new descriptor.
```

20

The *dup2* system call. Example 1

```
#include <fcntl.h>
...
#include <stdio.h>

int main() {
    int f1 = open("f1", O_WRONLY);
    dup2(f1, 1);
    dup2(1, 2);
    printf("Example dup2\n");
    perror("Example error\n");
    return 0;
}
```

Example: Test Program 2018

21

The *dup2* system call. Example 2

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <unistd.h>

int main() {
    int oldfd, newfd;

    oldfd = open("ERRORS", O_CREAT | O_WRONLY, 0755);
    if (oldfd < 0) {
        fprintf(stderr, "on stderr: open ERRORS impossible\n");
        fprintf(stdout, "on stdout: open ERRORS impossible\n");
        wait(1);
    }

    newfd = dup2(oldfd, 2); /*automatically closes the standard stderr file*/
    /*the new default error file will be ERRORS*/
}
```

22

The *fcntl* system call

```
#include <unistd.h>
#include <fcntl.h>

int fcntl(int fd, int cmd, ... /* arg */);

Provides or changes properties of an open file;

cmd:
  F_DUPFD - duplicating a descriptor
  F_GETFD (or F_SETFD) - file descriptor flags
  F_GETFL (or F_SETFL) - file status flags
  F_GETOWN (or F_SETOWN) - manipulation of I/O signals
  F_GETLK (or F_SETLK) - blocking files
```

The *fcntl* system call. Example

23

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <stdio.h>

// fcntl example
int main() {
    int fd;
    int flags = O_RDONLY | O_CREAT | O_WRONLY | S_IRWXCU;
    printf("fcntl example\n");
    fd = open("fcntl", flags, 0700);
    if (fd < 0) {
        perror("fcntl");
        return 1;
    }
    close(fd);
    return 0;
}
```

fcntl(2) - Linux man page

File Locks

24

- **advisory lock**
 - **mandatory lock**
- *exclusive lock* - when a single process has access to the file or portion of the file;
- *shared lock* - when the shared portion can be read simultaneously by several processes, but no process writes.

Advisory lock by *fcntl*

25

```
struct flock {
    ...
    short l_type; /* Lock type: F_RLOCK, F_WRLCK, F_UNLCK */
    short l_whence; /* Interpretation mode l_start: SEEK_SET, SEEK_CUR, SEEK_END */
    off_t l_start; /* Offset start region */
    off_t l_len; /* Number of bytes blocked/unlocked */
    pid_t l_pid; /* PID of the process that already blocked the region (F_GETLK) */
    ...
};
```

Advisory lock by *fcntl* (cont.)

26

```
#include <fcntl.h>

int fcntl(int fd, int cmd, struct flock* lock);
```

cmd: can take one of the values:

F_GETLK	- check whether the indicated resource is locked or not;
F_SETLK	- blocks the resource specified in the third parameter;
F_SETLKW	- blocks only the desired resource. If it's already blocked, it'll wait for its release.

The *lockf* system call

27

```
#include <unistd.h>

int lockf(int fd, int cmd, off_t len);
```

cmd: can take one of the values:

F_ULOCK	- unlock a blocked region;
F_LOCK	- blocking a region;
F_TLOCK	- test and lock if it is not already locked;
F_TEST	- testing a region whether it's locked or not.

The End

28