

1

Unix processes

Process = a running program that uses a system resource set (memory, processor, disk, network interface, etc.)

2

The image of a process in memory

- **User context** – the portion of the address space accessible during execution by that process in user mode;
- **Kernel context** – maintained by the kernel and accessible only through specific system calls.

3

The image of a process in memory

Diagram illustrating the memory layout of a process. It shows the Process table, Process regions vector, Region table, Table of virtual memory pages, Memory pages, User area, File descriptors table, Registry control (hardware), Kernel stack, Table of file descriptors open in the system, Table of inodes, and File system.

5

Entry in the process table

- process state
- pointer to the User area and the Process Memory Regions Table
- process size in memory
- PID, PPID
- UID, EUID, GID, EGID
- process priority (number between 1 and 39 used for scheduling)
- signals sent to the process
- time statistics (e.g., using the processor)
- process memory status (if the process image is in the main memory or in the swap memory)
- pointer to the next process in the process queue that is in the READY state
- event descriptors that occurred while the process was in SLEEP state

6

User area

- pointer to the entry in the process table that corresponds to this user area
- UID, EUID, GID, EGID - used to determine process access rights
- timers - time spent in user mode and in kernel mode
- vector of actions to handle signals
- process control terminal - the one from which it was launched
- return values and possible errors after making system calls
- the current directory and the root directory
- environmental variables
- possible kernel restrictions imposed on the process (e.g., process size in memory)
- file descriptor table (data about all open process files)
- umask - access rights mask for newly created files by this process

7

Creating a Process. System call `fork`

```
#include <unistd.h>

pid_t fork();
```

Return:

- 1 - if the operation could not be performed (`errno`);
- 0 - in child code;
- pid - in parent code, where `pid` is the process identifier of the newly created child.

Diagram illustrating the memory layout of a process. It shows the Process table, Process regions vector, Region table, Table of virtual memory pages, Memory pages, User area, File descriptors table, Registry control (hardware), Kernel stack, Table of file descriptors open in the system, Table of inodes, and File system.

9

Using `fork` - suggestions - 1

```
...
if ( pid=fork() ) < 0 ) {
    perror("Error");
    exit(1);
}

if(pid==0) {
    /* child code */
    ...
    exit(0);
}

/* parent code */
...
wait(&status);
```

10

Using `fork` - suggestions - 2

```
pid=fork();
switch(pid) {
    case -1:
        printf(stderr, "ERROR: %s\n", sys_errlist[errno]);
        exit(1);
        break;

    case 0:
        /* child code */
        break;

    default:
        /* parent code */
        break;
}
```

11

Using `fork` - suggestions - 3

```
if (fork() == 0) {
    /* child code */
}

else {
    /* parent code */
}
```

12

Execution of an external program. System calls `exec*`

```
#include <unistd.h>

int execl(const char *path, const char *arg, ...);
int execip(const char *file, const char *arg, ...);
int execl(const char *path, const char *arg, ..., char * const envp[]);
int execvp(const char *path, char *const argv[]);
int execvp(const char *file, char *const argv[]);
```

man 3 exec

13

`execlp()` example

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

main() {
    pid_t pid;
    if ( (pid=fork()) == -1 ) perror("fork error");
    else if (pid==0) {
        execlp("newShell", "newShell", NULL);
        printf("Return not expected. Must be an execlp error!\n");
    }
}
```

14

`exec()` example

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

main() {
    fork();
    printf("The process id is %d\n", getpid());
    execl("/bin/sh", "sh", "-l", 0);
    printf("This line is not printed!\n");
}
```

15

`exec*` examples

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    // execl("/bin/sh", "/bin/sh", "-l", NULL);
    // execip("sh", "sh", "-l", NULL);
    // execl("/bin/sh", "/bin/sh", "-l", "pl.c", "execl.c", "fork1.c", "m", NULL);
    execl("/bin/sh", "/bin/sh", "-l", "-c", "execl.c", "fork1.c", "m", NULL);
}
```

16

The concept of `fork`

- `fork` creates a child process that is a *clone* of the parent
- child has a (virtual) copy of the parent's virtual memory
- child is running the same program as the parent
- child *inherits* open file descriptors from the parent
- child begins life with the same register values as parent
- the child process may execute a different program in its context with a separate `exec` system call

17

System calls `wait()` and `waitpid()`

```
#include <sys/types.h>
#include <sys/wait.h>

pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);

Return:
-1 - error
>0 - the process ID of the terminated child
```

man 7 wait

18

`wait()` system call. Example

```
int main() {
    int i, saved_status;
    for (i=0; i<3; i++) {
        if (fork()==0) {
            // child process
            exit(i);
        }
        while (wait(&saved_status) != -1); // loop till all children terminate
    }
}
```

19

Other functions for working with processes

```
#include <sys/types.h>
#include <unistd.h>

pid_t getpid(); // - returns the PID of the current process
pid_t getppid(); // - returns the PID to the parent of the current process
uid_t getuid(); // - returns the user ID that launched the current process
gid_t getgid(); // - returns the user group identifier that launched the current process
```

man 3 getpid, man 3 getuid, man 3 getgid

20

Library function `system`

```
#include <stdlib.h>

int system(const char *command);
```

- launches a program on the disk, using `fork()` followed by `exec()` together with `waitpid()` in the parent.

man 3 system

21

`vfork()` system call

```
#include <sys/types.h>
#include <unistd.h>

pid_t vfork(void);
```

- creates a new process, just like `fork()` but does not fully copy the parent's address space. It is used in conjunction with `exec()`, and has the advantage that it does not consume the time required for copy operations that would be useless anyway if `exec()` is called immediately (however, the process will be overwritten with the program taken from the disk).

man 3 vfork



22

Exercise

Explain the effect of the following code sequence:

```
int i;

for (i = 1; i <= 10; i++) fork();
```



- Created – the process is created through a *fork* system call.
- ReadyMemory – the process is in queue in RAM and is ready for execution;
the kernel decides when it enters into execution.
- ReadySwapped – the process is ready to be executed, but swapped due to memory shortage;
to enter the execution state must first enter the **ReadyMemory** state.
- RunKernel – process instructions are executed in the kernel mode (eg system calls, interrupt handlers).
- RunUser – program instructions are executed.



- | | |
|---------------------|---|
| Preempted | <ul style="list-style-type: none"> – almost identical to ReadyMemory. After running a system call, the process must be passed to Runnable or ReadyMemory. If there is a process with a higher priority in ReadyMemory, the current process is passed to Preempted and the process with the higher priority goes to Runnable, otherwise the current process passes to Runnable. A process in Preempted moves to Runnable when it has the highest priority. |
| SleepMemory | <ul style="list-style-type: none"> – the process is in memory but can not be executed until there is an event that will get it out of this state (e.g., waits for a resource). |
| SleepSwapped | <ul style="list-style-type: none"> – the process is swapped on the disk due to memory shortage. It can not be executed until an event occurs that will get it out of this state. |
| Zombie | <ul style="list-style-type: none"> – the process no longer exists, has ended, but has not been removed from the process table. |

The End